

# Confidence Estimation for Trustworthy and Efficient Speech Systems

---

Part 1: Theory (85 min)

Break (10 min)

Part 2: Applications (85 min)

# Confidence Estimation for Trustworthy and Efficient Speech Systems

---

Part 1: Theory

Vipul Arora

# Outline

---

- Introduction
- Assessment
- Why mis-calibration happens?
- Confidence calibration: post hoc methods and Bayesian methods
- Disentangling sources of uncertainty: Epistemic and Aleatoric

# Future of AI

---

- Current AI



<https://voicebot.ai/>



Wikipedia.com



Wikipedia.com

- Future AI

- AGI

- Human-Machine Collaboration




Dreamstime.com

# Mistakes and Trust



Hello, I am  
giving the  
stock



Take  
paracetamol



Weather Tomorrow  
8 AM

$16^{\circ}\text{C} \pm 1.2^{\circ}\text{C}$

# Need


---

- **Self-driving cars**: obstruction or not? Defer to human driver
- **Healthcare**: Operate or not? Defer to human doctor
- **Finance**: invest money or not? Defer to human expert
- **Screening**: accept or not? Defer to human examiner

# For Classification

# Deep Networks

- Output class



$$y = \arg \max_j o_j$$

- Confidence,  $P(\text{output}=\text{correct})$




# How to assess the calibration?

Guo et al., On Calibration of Modern Neural Networks, ICML 2017

# Intuitively

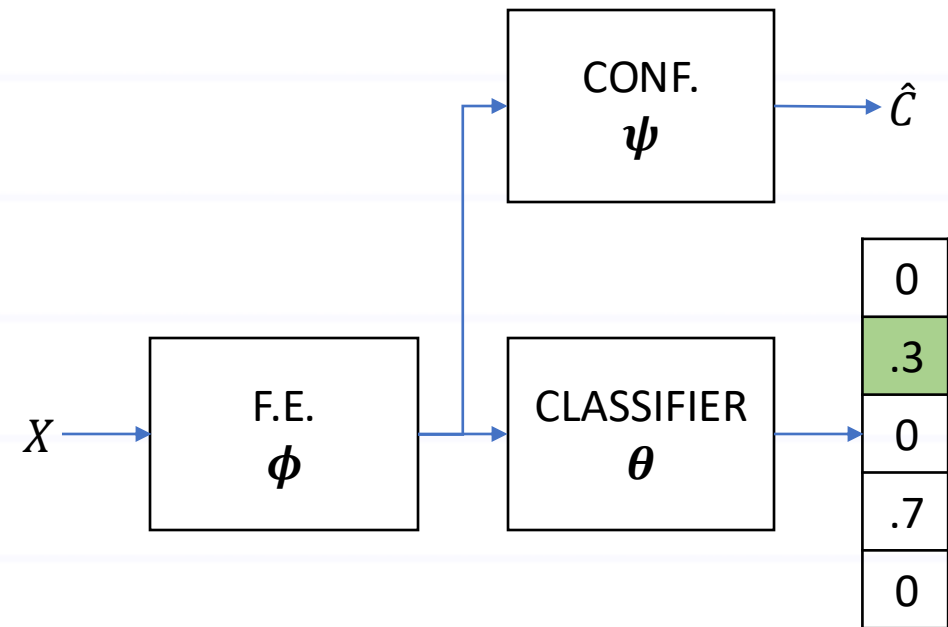
---

- Large confidence implies correct output
- Small confidence implies incorrect output



# Notations

- input:  $X$
- ground truth:  $Y \in \mathcal{Y} = \{1, \dots, K\}$
- output:  $\hat{Y} = f_{\theta} \left( f_{\phi}(X) \right)$
- output class:  $\hat{Y}^* = \operatorname{argmax}_k \hat{Y}[k]$
- confidence:  $\hat{C} = f_{\psi} \left( f_{\phi}(X) \right)$



# Calibration

---

- **Que:** When is the model perfectly calibrated,  $P(\hat{Y}^* = Y | \hat{C} = c) = ?$
- $P(\hat{Y}^* = Y | \hat{C} = c) = c, \forall c \in [0,1]$
- **Que:** what terms above are functions of  $X$ ?
- $Y, \hat{Y}^*, \hat{C}$
- **Que:** what terms above are independent variables?
- $c$
- **Que:** How do you compute  $P(\hat{Y}^* = Y | \hat{C} = c)$  from samples?
- = # of correct / total #, for all samples with  $\hat{C} = c$


# Problem

---

$$P(\hat{Y}^* = Y | \hat{C} = c) = c, \forall c \in [0,1]$$

- $c$  is a continuous variable. How many  $X$  in a finite database can have  $\hat{C} = c$
- So, we need approximation
- E.g., binning of  $c$

# Reliability Diagrams



- Accuracy vs confidence
- Let  $B_m$  be the set of samples with  $\hat{C}$  falling in  $m$ th bin
- Accuracy,  $\text{acc}(B_m) = \frac{1}{|B_m|} \sum_{i \in B_m} 1(\hat{y}_i^* = y_i)$
- Confidence,  $\text{conf}(B_m) = \frac{1}{|B_m|} \sum_{i \in B_m} \hat{C}_i$
- For perfectly calibrated model,  $\text{acc}(B_m) = \text{conf}(B_m) \forall m = \{1, \dots, M\}$

# Exercise

---

For a flower classification task,


$$Y = \{RRRR \text{ } JJJJ \text{ } LLLL\}$$

$$\hat{Y}^* = \{RLRR \text{ } JLJR \text{ } LLRJ\}$$

$$\hat{C} = \{.8, .7, .4, .7, \quad .7, .8, .8, .4, \quad .8, .7, .4, .4\}$$

Draw the reliability diagram.


# Expected Calibration Error



- Average over the reliability diagram
- $$\text{ECE} = \mathbb{E}_{\hat{C}} [ |P(\hat{Y}^* = Y | \hat{C} = c) - c| ]$$
- **Que:** write the Monte Carlo approximate of ECE
- $$\text{ECE} = \sum_{m=1}^M \frac{|B_m|}{n} |\text{acc}(B_m) - \text{conf}(B_m)|$$





# Maximum Calibration Error




- Take max error over the reliability diagram
- $$\text{MCE} = \max_{c \in [0,1]} |P(\hat{Y}^* = Y | \hat{C} = c) - c|$$
- **Que:** write the Monte Carlo approximate of MCE
- $$\text{MCE} = \max_{m \in \{1, \dots, M\}} |\text{acc}(B_m) - \text{conf}(B_m)|$$

# Real world systems

- Confidence = probability of being correct



## Calibrated



# Why are Deep Networks Mis-calibrated?

# Classification

---




# Feature Extraction

---


- height =  $x_1$
- diameter =  $x_2$



Que: confidence?



**Que:** What are roughly  $w_1$ ,  $w_2$ ,  $b$ ?





Hard  
Classification


■ = ★

Soft  
Classification

$$P(\blacksquare = \star) = 0.6$$


$$P(\blacksquare = \bullet) = 0.4$$

$$\hat{c} = 0.6$$



# Maths


- $z = w_2x_2 + w_1x_1 + b$
- **Que:** What is the output of a binary classifier?



- $P(y = 1|x) = \frac{1}{1+\exp(-z)} = \sigma(z)$
- **Que:** Keeping the class boundary same, what decides this  $P$ ?
- $P(y = 1|x) = \sigma(w \times (w_2x_2 + w_1x_1 + b))$

# Why mis-calibration

- **Que:** What is the loss function?
- $\text{Loss} = -\mathbb{E}[y \ln \hat{y} + (1 - y) \ln(1 - \hat{y})]$
- It is min if  $\hat{y} = y \in \{0,1\}$
- The loss does not favor fractional  $\hat{y}$
- $w$  is untuned



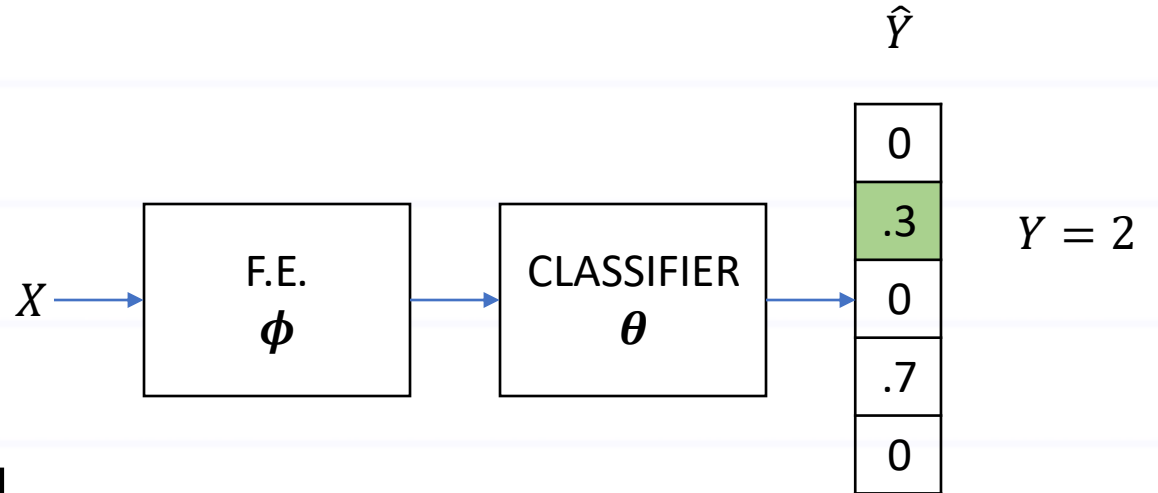
# Confidence Calibration

Guo et al., On Calibration of Modern Neural Networks, ICML 2017


# Goal

- Derive  $\hat{C}$  using  $\hat{Y}, \hat{Y}^*, Z, X$

$$\hat{C}(\hat{Y}, \hat{Y}^*, Z, X)$$

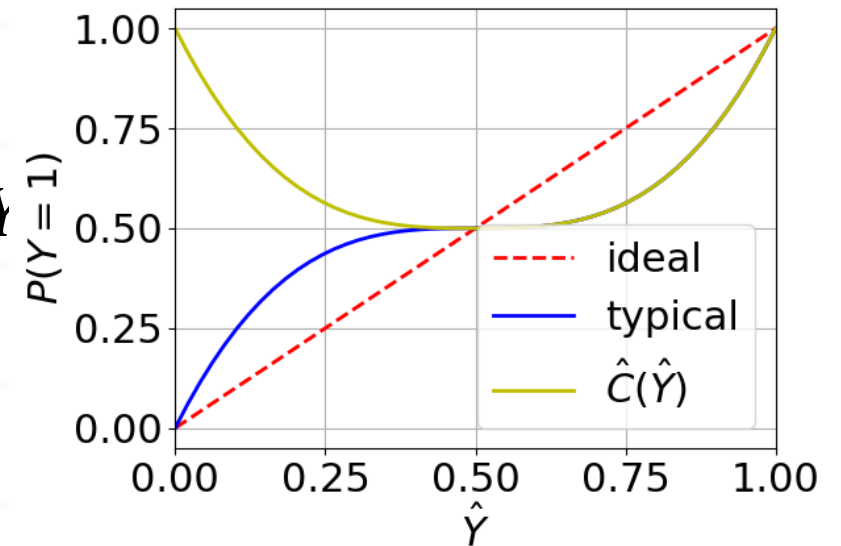


- Can't use  $Y$  during testing
- Let's focus on binary classification
- Que:** Draw  $\hat{Y}$  vs  $Z$




# Estimating $\hat{C}(\hat{Y})$

- **Que:** Draw typical  $P(Y = 1)$  vs  $\hat{Y}$
- **Que:** What should be  $\hat{C}(\hat{Y})$  if you know  $P(Y = 1)$  curve?
- **Ans:**  $\hat{C}(\hat{Y}) = P\left(Y = \arg \max_k \hat{Y}\right)$ . Can you draw it for binary case?
- BTW, Instead of modeling  $\hat{C}(\hat{Y})$ , we could model  $\hat{C}(k) := P(Y = k)$  for all  $k$




# 1. Histogram Binning Method

- to estimate  $\hat{C}(\hat{Y})$
- Divide  $\hat{Y} \in [0,1]$  into  $M$  bins
- Let  $P(Y = 1) = \theta_m$  if  $\hat{Y}$  falls in bin  $m$
- To estimate  $\theta_m$ ,  $\text{Loss} = \sum_m \sum_i 1(\hat{Y}_i \in \text{bin } m) (\theta_m - Y_i)^2$
- **Que:** What is the optimal  $\theta_m$ ?
- $$\theta_m = \frac{\sum_i 1(\hat{Y}_i \in \text{bin } m) Y_i}{\sum_i 1(\hat{Y}_i \in \text{bin } m)}$$



## 2. Isotonic Regression Method

- to estimate  $\hat{C}(\hat{Y})$
- One can learn  $P(Y = 1)$  as a function of  $\hat{Y}$  using simple regression models
- Isotonic regression model is one such model






### 3. Platt Scaling Method

- to estimate  $\hat{C}(Z)$   $Z$  is logit
- **Que:** Draw typical  $P(Y = 1)$  vs  $Z$
- Approximate this with a sigmoid as

$$P(Y = 1) = \sigma(aZ + b)$$

- If  $b = 0$ , it is called Temperature Scaling method with  
 $a = 1/T$



### 3. Platt Scaling Method

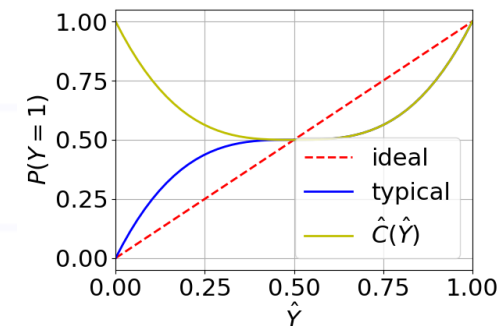
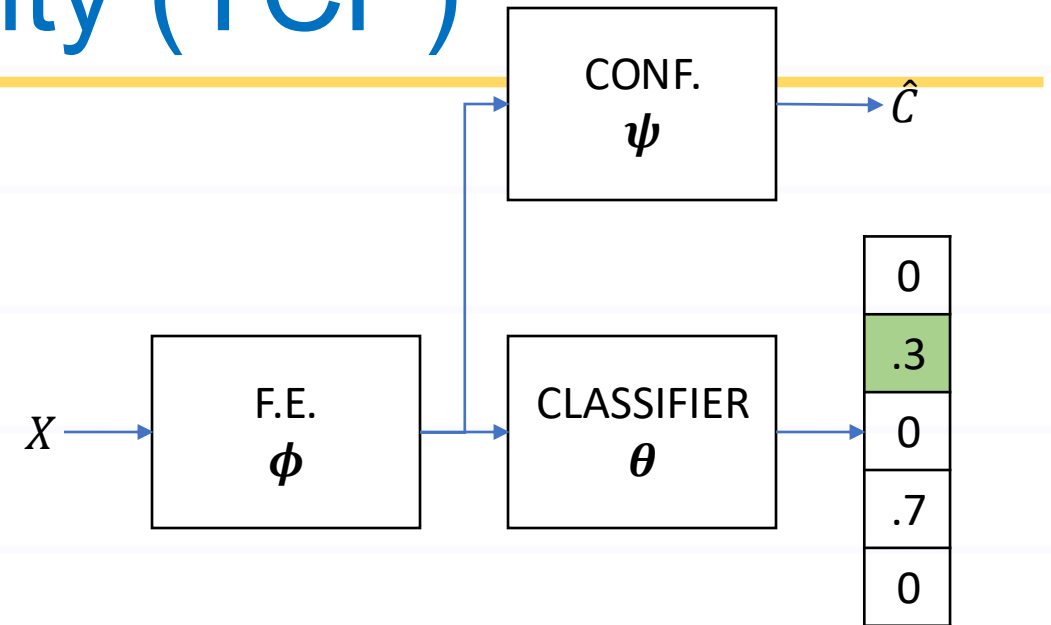
---

$$P(Y = 1) = \sigma(aZ + b)$$

- $a$  and  $b$  are estimated using MLE over validation data
- **Que:** Could you write the loss function?
- Loss =  $-\sum_i \delta_{y_i,1} \ln \sigma(aZ + b) + \delta_{y_i,0} \ln(1 - \sigma(aZ + b))$

## 4. True Class Probability (TCP)

- $\hat{C}(X)$
- Train  $\psi$  as a regressor
- $\text{Loss} = \mathbb{E} \left[ \left( C(X) - \hat{C}(X) \right)^2 \right]$
- What should be the target?
- Use  $C(X) = \hat{Y}[k] ; k = Y$

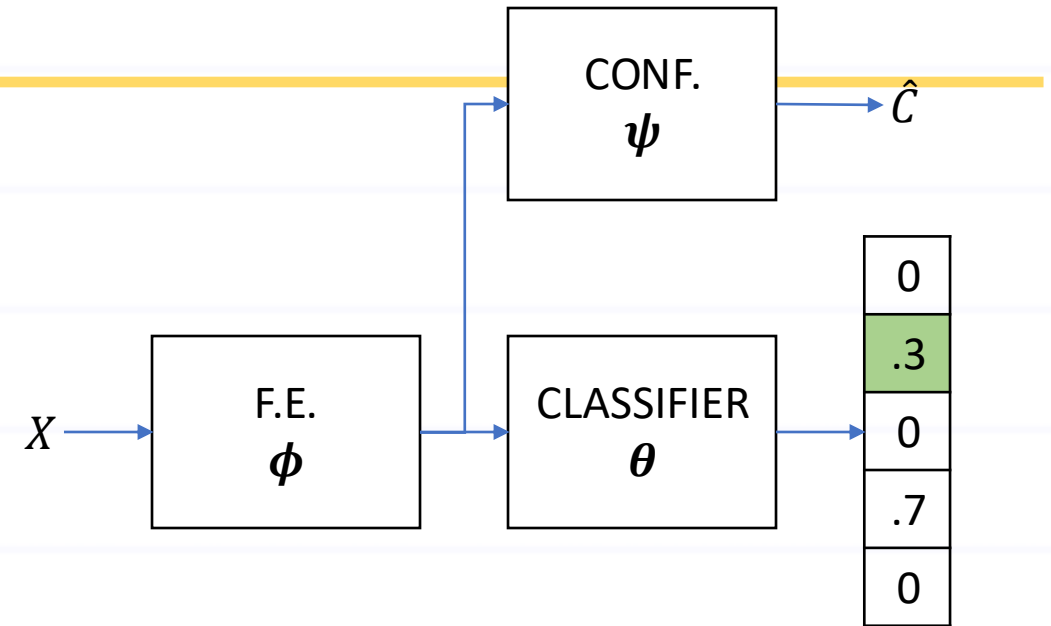


Corbiere, et al., "Addressing failure prediction by learning model confidence," NeurIPS 2019.

## 4. Normalized TCP

- $\hat{C}(X)$

- $\text{Loss} = \mathbb{E} \left[ \left( C(X) - \hat{C}(X) \right)^2 \right]$



- When number of classes is large,  $\hat{Y}[k]$  gets smaller


- Use  $C(X) = \frac{\hat{Y}[k]}{\max_{k'} \hat{Y}[k']} ; k = Y$

Corbiere, et al., “Addressing failure prediction by learning model confidence,” NeurIPS 2019.

# For multi-class classification

---

- Treat it as  $K$  one-vs-all problems
- Estimate  $P(Y = k)$  for all  $k$  using  $Z[k]$  or  $Z$
- Hence, get  $\hat{C}[k]$
- **Que:** Why not treat it as a full multi-class problem?
- Curse of dimensionality




# Bayesian Methods

Gal and Ghahramani, Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning, ICML 2016

# Bayesian Neural Network

- Typical NN

$$\hat{p}(Y|X) = \hat{Y} = f_{\theta} \left( f_{\phi}(X) \right)$$



- In Bayesian NN,  $\phi, \theta$  are also random variables; so, we get

$$p(Y|X) = \iint p(Y|X, \theta, \phi) p(\theta, \phi) d\theta d\phi$$

# Monte Carlo Estimation


---

- $p(Y|X) = \iint p(Y|X, \theta, \phi) p(\theta, \phi) d\theta d\phi$
- $p(Y|X) = \mathbb{E}_{\theta, \phi \sim p(\theta, \phi)} [p(Y|X, \theta, \phi)]$
- $\hat{p}(Y|X) \approx \frac{1}{N} \sum_i p(Y|X, \theta_i, \phi_i); \quad \theta_i, \phi_i \sim p(\theta, \phi)$



# Monte Carlo Dropouts

- $\hat{p}(Y|X) \approx \frac{1}{N} \sum_{i=1}^N p(Y|X, \theta_i, \phi_i) ; \theta_i, \phi_i \sim p(\theta, \phi)$
- $\theta, \phi \sim p(\theta, \phi)$  is approximated using dropout



# Uncertainty

---

- $\mu = \frac{1}{N} \sum_i \hat{Y}_i$
- $\Sigma = \frac{1}{N} \sum_i (\hat{Y}_i - \mu)^\top (\hat{Y}_i - \mu)$
- Indicators of uncertainty:
  - Total variance =  $\sum_{k,l} \Sigma_{kl}$
  - Entropy =  $-\sum_k \mu[k] \ln \mu[k]$

# Ensemble Method

---

- $\hat{p}(Y|X) \approx \frac{1}{N} \sum_i p(Y|X, \theta_i, \phi_i) ; \quad \theta_i, \phi_i \sim p(\theta, \phi)$
- $\theta, \phi \sim p(\theta, \phi)$  is a model trained using stochastic optimization algorithms
- We train  $N$  different models and use them to estimate uncertainty

# Input Perturbation Method

---

- $p(Y|X) = \int p(Y|\tilde{X})p(\tilde{X}|X)d\tilde{X}$
- $\hat{p}(Y|X) \approx \frac{1}{N} \sum_i p(Y|\tilde{X}_i); \quad \tilde{X}_i = X + \epsilon_i$
- Here,  $\epsilon_i$  is noise or systematic perturbation of input  $X$


# Regression

Kuleshov et al., Accurate Uncertainties for Deep Learning Using Calibrated Regression, ICML 2018

# Regression


---

- Output  $\hat{Y} \in \mathbb{R}^d$
- How do we define confidence?



# Confidence

- Let us estimate the pdf of output, instead of point estimate. Let  $y$  be the output random variable.
- Now, true value  $Y$  can lie anywhere




# Confidence

- Example: Let  $y_2$  be such that

$$\int_{-\infty}^{y_2} p(y|X) dy = 0.5$$

- Now, empirically  $P(Y \leq y_2|X)$  should be 0.5
- This is true for only one  $X$ , how to generalize







# Confidence

- Let us define an interval in terms of probability mass
- Example: Let  $y_2$  be a point such that

$$\int_{-\infty}^{y_2} p(y|X) dy = 0.5$$

- But we know that this is  $F_Y(y_2|X) = 0.5$ , where  $F_Y(y_2)$  is the distribution function or CDF
- Thus,  $y_2 = F_Y^{-1}(0.5|X)$
- $y_2$  is a function of  $X$



# Confidence

---

- $y_2 = F_Y^{-1}(0.5|X)$
- Now, for a calibrated regression model, empirically


$$\frac{1}{N} \sum_n^N \mathbb{I}\{Y_i \leq F_Y^{-1}(0.5|X)\} = ?$$

- $= 0.5$

# Confidence

- In general, for a calibrated regression model,

$$\frac{1}{N} \sum_n^N \mathbb{I}\{Y_i \leq F_Y^{-1}(c|X)\} = c$$



# Confidence

---

- Even more general, for a calibrated regression model,

$$\frac{1}{N} \sum_n^N \mathbb{I}\{F_Y^{-1}(c_2|X) \leq Y_i \leq F_Y^{-1}(c_2|X)\} = c_2 - c_1$$

- $(c_2 - c_1)$  is the confidence interval

# Summary

---

- Confidence calibration is possible when we estimate the complete pdf  $p(Y|X)$  instead of just a point estimate  $\hat{Y}$
- We need to estimate CDF  $F_Y(y|X)$  and inverse CDF  $F_Y^{-1}(c|X)$  to get confidence intervals

# Confidence

---

- It is common to assume output to be Gaussian

$$p(y|X) = \mathcal{N}(y; \mu(X), \sigma^2(x))$$

- $\mu(X), \sigma(x)$  are estimated using NNs





# Sources of Uncertainty

- Kendall and Gal, What Uncertainties Do We Need in Bayesian Deep Learning for Computer Vision?, NeurIPS 2017
- Hüllermeier and Waegeman, Aleatoric and epistemic uncertainty in machine learning: an introduction to concepts and methods, Machine Learning 2021

# Calibration is not enough

Two kinds of uncertainty:

1. **Model** is limited, not trained on this data or class. **Epistemic Uncertainty**
2. **Data** ambiguous, even if model has been trained on similar data. **Aleatoric Uncertainty**





# Should we distinguish?

---

- Epistemic Uncertainty
  - tells about out of domain data (new, unseen inputs) (useful for model adaptation and active learning)
  - tells about anomalies and outliers
  - tells about unseen classes (novel class discovery)
- Aleatoric Uncertainty
  - tells about difficult data which needs manual intervention. More training won't help.

# Can we distinguish?


---

Yes

- Bayesian NN
- Evidential learning [Sensoy et al., Evidential Deep Learning to Quantify Classification Uncertainty, NeurIPS 2018]

# Consider the $Z$ space (logits)

- Assume Gaussian output




$$Z \sim \mathcal{N}(\mu_z, \sigma_z^2); \mu_z \in \mathbb{R}^K, \sigma_z \in \mathbb{R}^K$$

where  $\mu_z, \sigma_z = f_\theta \left( f_\phi(X) \right)$

- Output is as usual  $\hat{Y} = \text{softmax}(Z)$  for classification and  $\hat{Y} = Z$  for regression


# Uncertainty



- $\sigma_Z$  quantifies the uncertainty in  $Z$ , but what kind of uncertainty?
- Let us perturb the model parameters
- $\mu_{Z,i}, \sigma_{Z,i} = f_{\theta_i} \left( f_{\phi_i}(X) \right)$

# Uncertainty

- Que: What is  $\mathbb{E}[Z]$ ?



$$\mathbb{E}[Z] = \iint Z p(Z|X, w) p(w) dw dZ ; \quad w = \{\phi, \theta\}$$

$$= \int \mu_{z,w} p(w) dw$$

$$\approx \frac{1}{N} \sum_{i=1}^N \mu_{z,i}$$

# Uncertainty

---

- Que: What is  $\text{covar}[Z]$  or  $\mathbb{E}[(Z - \mathbb{E}[Z])^\top (Z - \mathbb{E}[Z])]$ ?

$$\mathbb{E}[Z^\top Z] = \iint Z^\top Z p(Z|X, w) p(w) dw dZ$$

$$= \int (\sigma_{Z,w}^2 I + \mu_{Z,w}^\top \mu_{Z,w}) p(w) dw$$

$$\approx \frac{1}{N} \sum_{i=1}^N (\sigma_{Z,i}^2 I + \mu_{Z,i}^\top \mu_{Z,i})$$

# Uncertainty


- $\text{covar}[Z] \approx \frac{1}{N} \sum_{i=1}^N (\sigma_{z,i}^2 I + \mu_{z,i}^\top \mu_{z,i}) - \mathbb{E}[Z]^2$

- $= \underbrace{\frac{1}{N} \sum_i \sigma_{z,i}^2 I}_{\text{Aleatoric Uncertainty}} + \underbrace{\frac{1}{N} \sum_i \mu_{z,i}^\top \mu_{z,i} - \left( \frac{1}{N} \sum_{i=1}^N \mu_{z,i} \right)^2}_{\text{Epistemic Uncertainty}}$

**Aleatoric  
Uncertainty**

**Epistemic  
Uncertainty**

**PICTORIAL  
UNDERSTANDING**



# Quick Introductions



# Evidential Learning

---

- Treat  $\mu_z, \sigma_z$  also as random variables
- $\text{var}[\mu_z]$  gives epistemic uncertainty
- $\mathbb{E}[\sigma_z^2]$  gives aleatoric uncertainty

# Conformal Prediction

---



- Random samples from an unknown distribution are given
- What is the probability that the next sample falls between the 3<sup>rd</sup> and 4<sup>th</sup> samples on the line?
- Hint: draw the CDF
- This gives us the confidence intervals

# Further Reading

---

- Lakshminarayanan et al., “*Simple and Scalable Predictive Uncertainty Estimation using Deep Ensembles*,” NeurIPS, 2017
- Seitzer, et al., “*On the Pitfalls of Heteroscedastic Uncertainty Estimation with Probabilistic Neural Networks*,” ICLR 2022
- Sensoy et al., “*Evidential Deep Learning to Quantify Classification Uncertainty*”, NeurIPS 2018
- Angelopoulos and Bates, “*A Gentle Introduction to Conformal Prediction and Distribution-Free Uncertainty Quantification*”, 2022
- Ryan Tibshirani, “*Conformal Prediction*”, Advanced Topics in Statistical Learning, Spring 2023

# Further Reading

---

- Nagarathna, Thishyan, Chaganti, Arora, “*ASR Confidence Estimation using True Class Lexical Similarity Score*”, Interspeech 2025
- Nagarathna, Thishyan and Arora, “*TeLeS: Temporal Lexeme Similarity Score to Estimate Confidence in End-to-End ASR*”, IEEE TASLP 2024
- Saxena and Arora, “*Interactive Singing Melody Extraction Based on Active Adaptation*”, IEEE TASLP 2024

# Questions?

## Next: Part 2